

# Taster: *Self-Tuning, Elastic and Online* Approximate Query Processing

Matthaios Olma

Odysseas Papapetrou

Raja Appuswamy

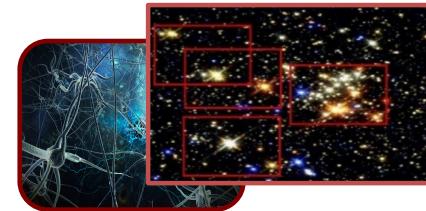
Anastasia Ailamaki



# Challenges of interactive data exploration

## Exploratory Applications

- Dynamic & data-driven



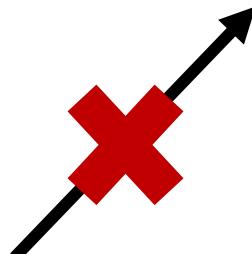
Scientific exploration

“Internet of Things” analytics

## Interactive response time



Reduce result precision  
– use AQP



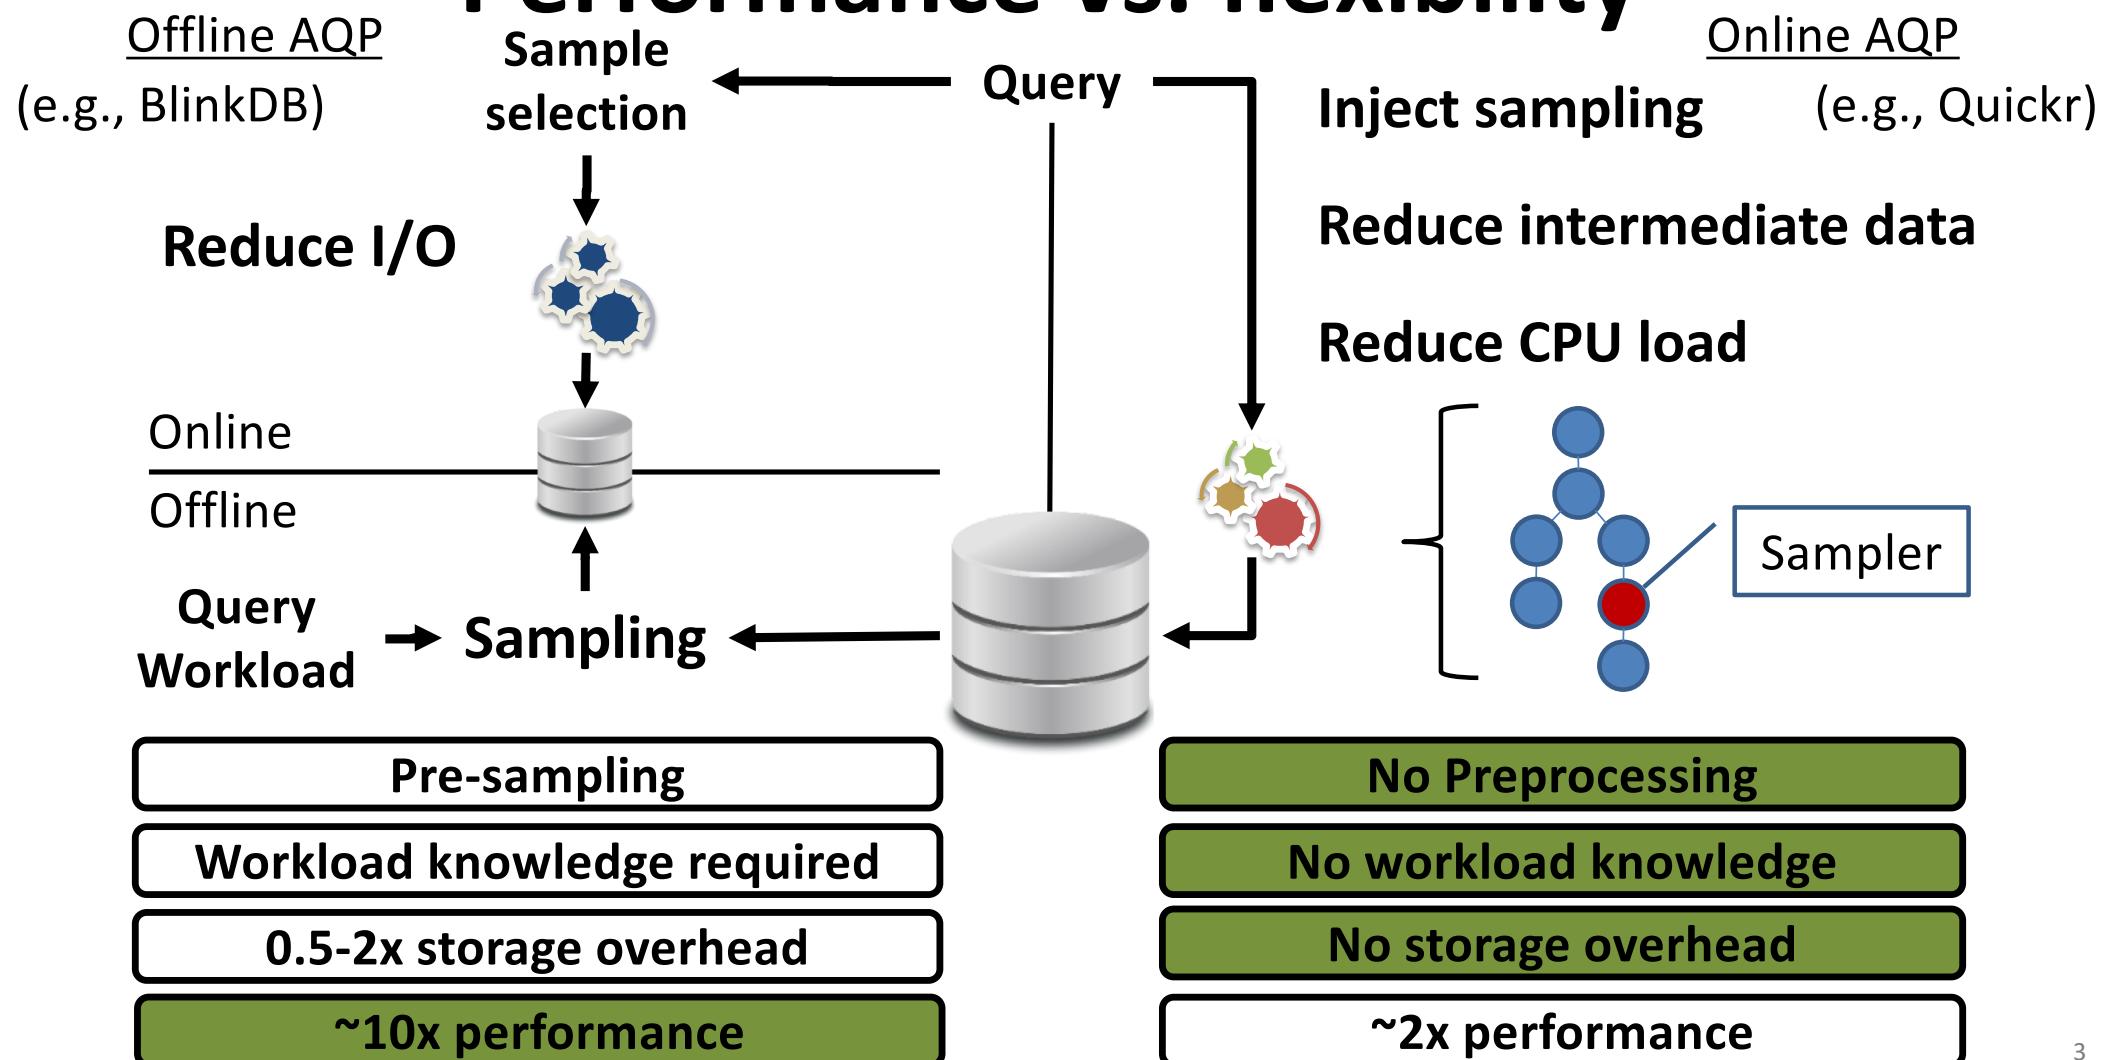
## Instant access to data



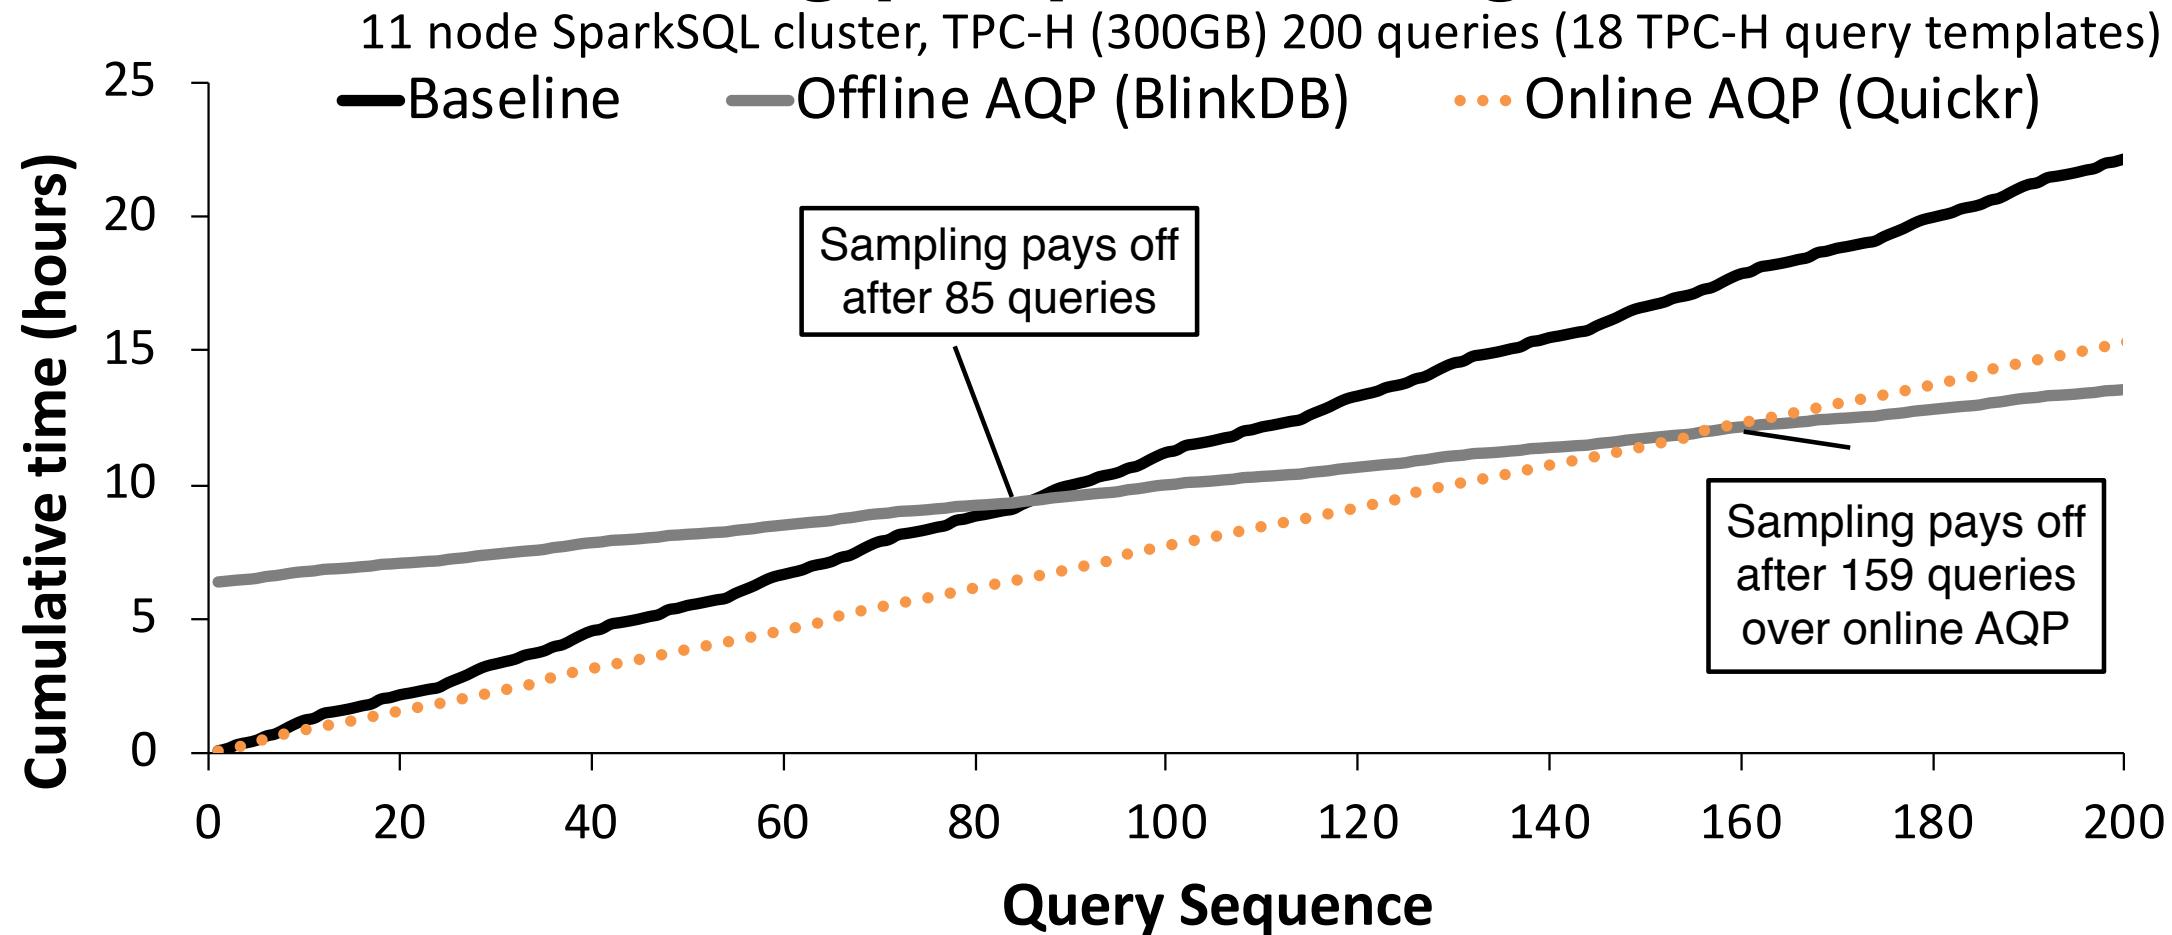
Building data  
summaries is expensive

Enable AQP with minimal pre-processing

# Performance vs. flexibility



# Reducing pre-processing time



Ideal: No sampling preparation cost & interactive access

# Enhancing Online Approx. Query Processing

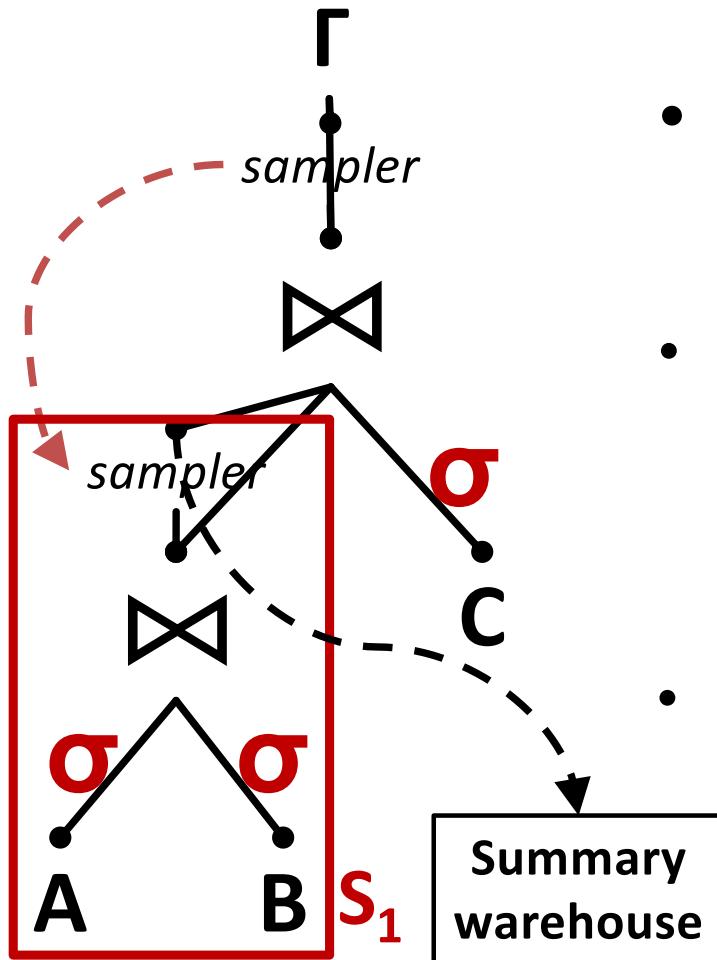
- Reduce the amount of data accessed
  - Materialize and Re-use intermediate generated summaries
- Adapt materialized summaries to workload and storage budget
- Use a variety of summaries other than samples

**What to  
materialize**

**If/When to  
materialize**

**If/When to  
evict**

# Materialize and re-use synopses



- Store all subplans and statistics in hitmap
  - Update when subplans re-appear
- Calculate prospective gains (cost:benefit)
  - Performance gains over future workload
  - Storage cost
  - Maximize benefit – Knapsack constraint problem
- Decide to materialize
  - Inject materializer operator
  - Store intermediate result in-memory and flush offline

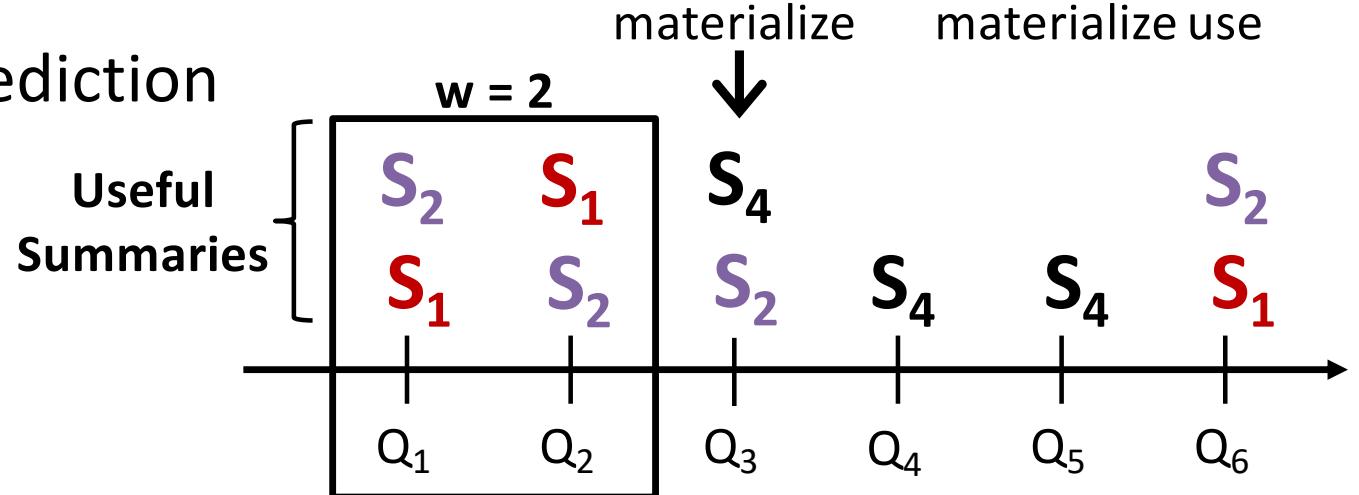
# Adapting materialized summaries

- Window-based prediction

Summary warehouse



Useful Summaries



- Ideal window size depends on: user, task, data
  - Keep statistics for  $(1-a)w$ ,  $w$ ,  $(1+a)w$
  - Adapt window size based on quality of predictions

**Abide to storage requirements despite workload shifts**

**Online tuning of window size improves forecast efficiency**

# Combining different data summaries

## Sampling

### All queries on subset of data

- Keep schema of original table
- Precision depends on query
- Uniform/Stratified sampling

Answer large subset of queries

Large size ~ 10% of input

I/O cost depending on size

## Sketches

### Some queries on all data

- Count/Sum/Avg
- Aggregations
- Single grouping attribute

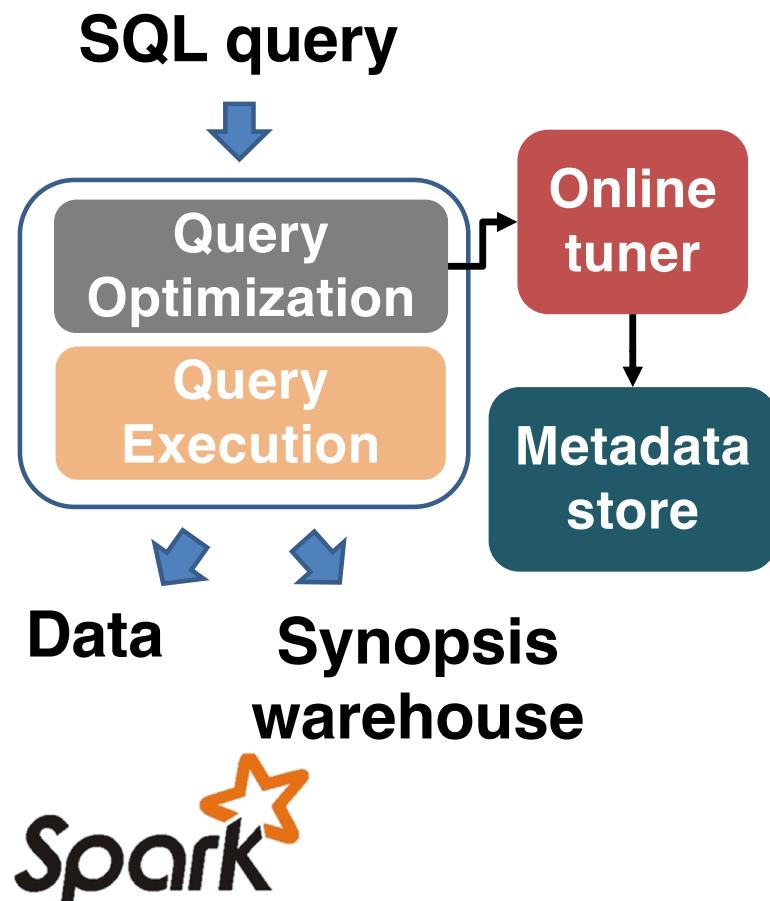
Answer specific queries

Compact ~KB

Constant access time

**Utilize each summary where useful**

# Taster Architecture



- Inject approximation operators into plans
- Re-use existing materialized synopses
- Choose which synopsis to generate
- Store statistics about the historical plans
- Store the synopsis over HDFS

# Experimental Setup

## Datasets

- TPC-H: sf300 (300GB), 18 query templates

## Systems

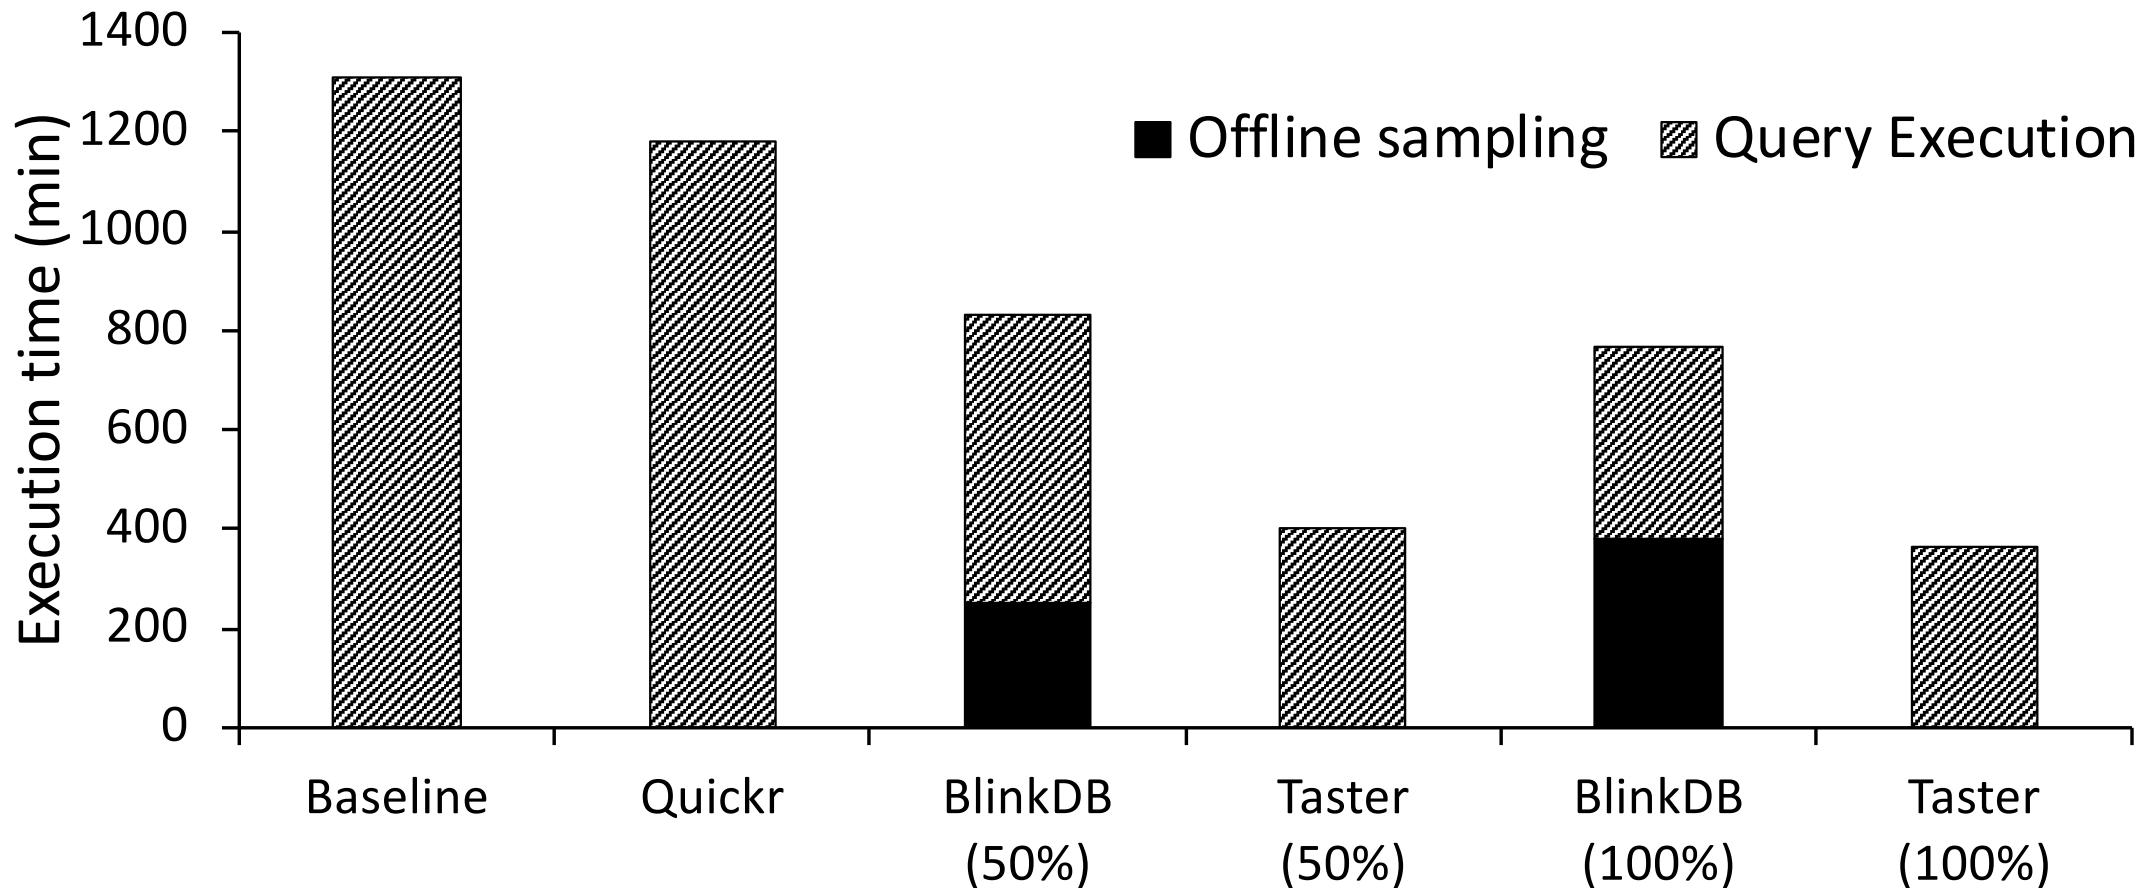
- SparkSQL (2.1.0)
- BlinkDB, Quickr, Taster over SparkSQL (2.1.0)

## Hardware

- 11 nodes x 2 x Intel Xeon X5660 CPU @ 2.80GHz, 48GB RAM, 10GbE (fix for each)

# End-to-End execution time

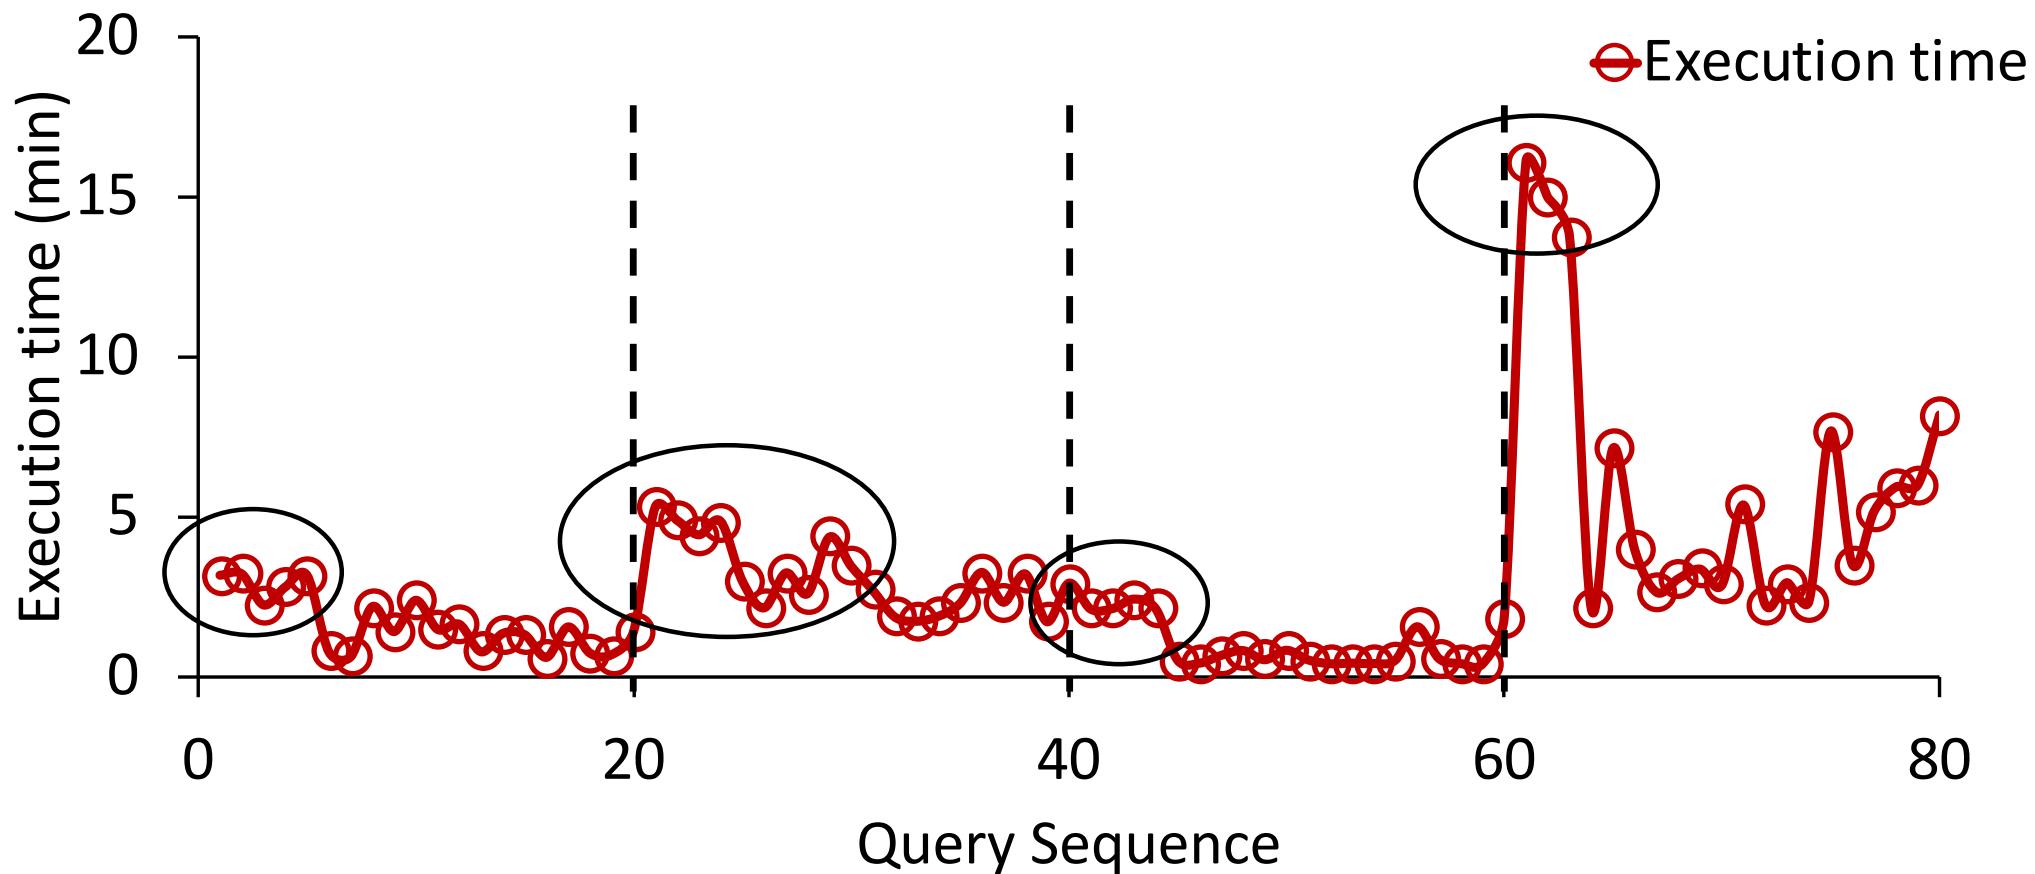
11 node SparkSQL cluster, TPC-H sf300, 200 queries (18 TPC-H templates)



**Taster offers comparable execution time to state-of-the-art**

# Adapting to shifting workload

11 node SparkSQL cluster, TPC-H sf300, 80 queries (18 TPC-H templates)



**Taster adapts efficiently to changes in workload**

# Take home message

- Piggy-back the creation of summaries over the query execution
  - In the context of distributed approximate query processing
- Adapt data summaries to workload shifts and reduce storage budget
- Provide query performance comparable to offline AQP approaches
  - With reduced building and storage cost

Thank you!